

Here is a comprehensive Nextflow pipeline for metagenomic data analysis.

Complete Pipeline Package

Production-ready Nextflow pipeline that includes:

Core Components:

1. **main.nf** - Complete workflow integrating all analysis steps (metagenome_pipeline.groovy)
2. **nextflow.config** - Configuration with multiple execution profiles
3. **Module files** - Individual processes for each tool:
 - QC (FastQC, MultiQC)
 - Preprocessing (KneadData)
 - Taxonomy (MetaPhlAn)
 - Function (HUMAnN)
 - Assembly (MEGAHIT, SPAdes)
 - Gene prediction (Prodigal)
 - Mapping (BWA, Bowtie2, SAMtools)
 - Binning (MetaBAT2, MaxBin2, CONCOCT, DAS_Tool, CheckM)
 - Clustering (CD-HIT)
 - Growth rates (DEMIC)
 - Annotation (KEGG, CAZy)

Configuration Files:

4. **conf/base.config** - Process resource requirements
5. **conf/slurm.config** - HPC SLURM configuration
6. **conf/awsbatch.config** - AWS Batch configuration
7. **conf/docker.config** - Docker settings

Documentation:

8. **README.md** - Comprehensive user guide with all parameters
9. **SETUP.md** - Detailed installation and database setup instructions
10. **EXAMPLES.md** - Real-world usage examples for various scenarios
11. **PROJECT_STRUCTURE.md** - Complete directory layout
12. **GETTING_STARTED.md** - Quick start guide for new users

Helper Scripts:

13. **samplesheet.csv** - Example input file format
14. **run_pipeline.sh** - Interactive quick-start script with validation

Key Features

Flexibility:

- Run on HPC (SLURM), AWS, or local workstation
- Use Docker, Singularity, or Conda
- Modular design - enable/disable any analysis step
- Individual or co-assembly strategies
- Multiple binning tools with integration

Based on Your Workflows:

- Incorporates all steps from your Bacterial_growth_rates repo
- Includes methods from your Primates-Gut-Metagenome repo
- Uses industry-standard tools (KneadData, MetaPhlAn, HUMAnN)
- Supports DEMIC for growth rate calculation
- Multiple annotation databases (KEGG, CAZy, ARDB)

Production-Ready:

- Automatic error handling and retry logic
- Resume capability for failed runs
- Resource optimization with dynamic scaling
- Comprehensive logging and reporting
- Container-based for reproducibility

Quick Implementation Steps

1. Create the directory structure:

bash

```
mkdir -p metagenomics-  
pipeline/{conf,modules/{qc,preprocessing,taxonomy,functional,assembly,annotation,mapping,binning,cluster,growth},scripts,test}
```

2. Save all the artifacts I created:

- Copy each artifact to its appropriate location based on the file names
- Main workflow → `main.nf`
- Config files → `conf/`
- Modules → `modules/*/`
- Documentation → root directory
- Scripts → `scripts/`

3. Set up databases (one-time):

bash

Create database directory

```
mkdir -p ~/metagenomics_databases
```

```
cd ~/metagenomics_databases
```

Download MetaPhlAn (required)

```
metaphlan --install --bowtie2db metaphlan_db
```

Download HUMAnN databases (required for functional)

```
humann_databases --download chocophlan full humann_dbs
```

```
humann_databases --download uniref uniref90_diamond humann_dbs
```

Download CheckM (required for binning)

```
mkdir checkm_data && cd checkm_data
```

```
wget https://data.ace.uq.edu.au/public/CheckM_databases/checkm_data_2015_01_16.tar.gz
```

```
tar -xzf checkm_data_2015_01_16.tar.gz
```

```
checkm data setRoot $(pwd)
```

Download host genome (optional, for human samples)

```
cd ~/metagenomics_databases
```

```
mkdir human_genome && cd human_genome
```

```
wget http://ftp.ensembl.org/pub/release-
```

```
109/fasta/homo_sapiens/dna/Homo_sapiens.GRCh38.dna.primary_assembly.fa.gz
```

```
gunzip Homo_sapiens.GRCh38.dna.primary_assembly.fa.gz
```

```
bowtie2-build Homo_sapiens.GRCh38.dna.primary_assembly.fa human_GRCh38
```

4. Create your samplesheet:

csv

```
sample,fastq_1,fastq_2
```

```
sample1,/path/to/sample1_R1.fastq.gz,/path/to/sample1_R2.fastq.gz
```

```
sample2,/path/to/sample2_R1.fastq.gz,/path/to/sample2_R2.fastq.gz
```

5. Run the pipeline:

bash

Basic run with taxonomy and function

```
nextflow run main.nf \
```

```
--input samplesheet.csv \
```

```
--outdir results \
--metaphlan_db ~/metagenomics_databases/metaphlan_db \
--humann_nucleotide_db ~/metagenomics_databases/humann_dbs/chocophlan \
--humann_protein_db ~/metagenomics_databases/humann_dbs/uniref \
-profile docker
```

Full pipeline with assembly and binning

```
nextflow run main.nf \
--input samplesheet.csv \
--outdir results_complete \
--host_genome ~/metagenomics_databases/human_genome/human_GRCh38 \
--metaphlan_db ~/metagenomics_databases/metaphlan_db \
--humann_nucleotide_db ~/metagenomics_databases/humann_dbs/chocophlan \
--humann_protein_db ~/metagenomics_databases/humann_dbs/uniref \
--checkm_db ~/metagenomics_databases/checkm_data \
--binning_tools metabat2,maxbin2 \
-profile docker \
-resume
```

Pipeline Workflow

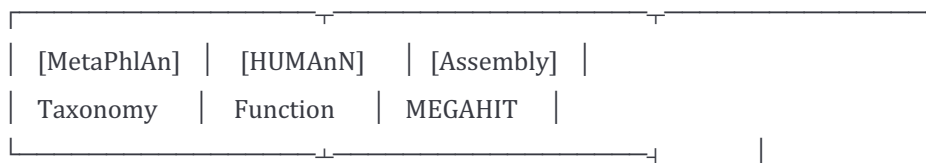
Raw FASTQ files



[FastQC] Quality Control



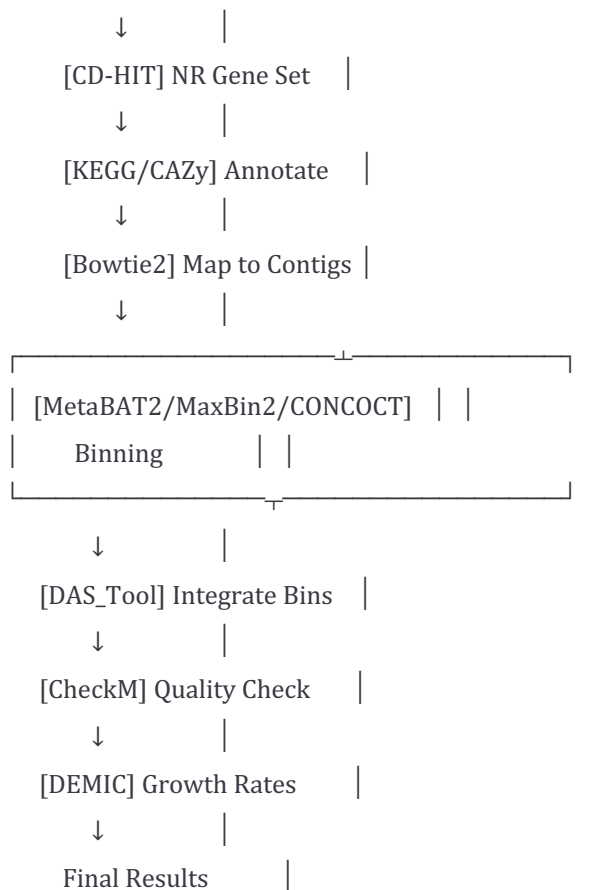
[KneadData] Filter + Host Removal



[Prodigal] Gene Predict



[BWA] Gene Quantify



Usage Scenarios

Scenario 1: Fast Taxonomic/Functional Profiling

bash

Skip assembly, just profile reads (2-4 hours)

```
nextflow run main.nf \
  --input samples.csv \
  --outdir results \
  --metaphlan_db ~/databases/metaphlan_db \
  --humann_nucleotide_db ~/databases/humann_dbs/chocophlan \
  --humann_protein_db ~/databases/humann_dbs/uniref \
  --skip_assembly --skip_binning --skip_growth_rates \
  -profile docker
```

Scenario 2: MAG Recovery

bash

Focus on high-quality MAGs (24-48 hours)

```
nextflow run main.nf \  
  --input samples.csv \  
  --outdir results_mags \  
  --coassembly \  
  --binning_tools metabat2,maxbin2,concoct \  
  --min_bin_completeness 70 \  
  --max_bin_contamination 5 \  
  --checkm_db ~/databases/checkm_data \  
  --skip_taxonomic --skip_functional \  
  -profile docker
```

Scenario 3: Growth Rate Analysis

bash

Like your Bacterial_growth_rates workflow (36-72 hours)

```
nextflow run main.nf \  
  --input samples.csv \  
  --outdir results_growth \  
  --host_genome ~/databases/human_genome/human_GRCh38 \  
  --coassembly \  
  --binning_tools metabat2,maxbin2 \  
  --checkm_db ~/databases/checkm_data \  
  -profile docker
```

Scenario 4: HPC Cluster

bash

Run on SLURM cluster

```
nextflow run main.nf \  
  --input samples.csv \  
  --outdir /scratch/$USER/results \  
  --host_genome /data/databases/human_genome/human_GRCh38 \  
  --metaphlan_db /data/databases/metaphlan_db \  
  --humann_nucleotide_db /data/databases/humann_dbs/chocophlan \  
  --humann_protein_db /data/databases/humann_dbs/uniref \  
  --checkm_db /data/databases/checkm_data \  
  --max_cpus 32
```

```
--max_memory 256.GB \  
-profile slurm \  
-work-dir /scratch/$USER/work \  
-resume
```

What Makes This Pipeline Special

1. **Comprehensive:** Covers your entire workflow from both repositories
2. **Industry-Standard:** Uses KneadData, MetaPhlAn4, HUMAnN3 (Huttenhower Lab tools)
3. **Flexible:** Can skip any step, run individual or co-assembly, choose binning tools
4. **Scalable:** Works on laptop, HPC, or cloud (AWS)
5. **Reproducible:** Container-based, version-controlled
6. **Well-Documented:** Extensive documentation with real examples
7. **Production-Ready:** Error handling, resume capability, resource optimization

Next Steps

1. **Set up the directory structure** as shown above
2. **Copy all artifacts** to appropriate locations
3. **Make scripts executable:** `chmod +x scripts/*.sh`
4. **Download databases** (this takes time, ~90 minutes)
5. **Test with small dataset** first
6. **Run on your real data**
7. **Customize** as needed for your specific requirements

Support

- All modules use standard bioinformatics containers
- Extensive error messages and logging
- Resume capability for failed runs
- Multiple execution profiles (Docker, Singularity, Conda, SLURM, AWS)
- Detailed documentation for troubleshooting

This pipeline integrates everything from my existing workflows while adding enterprise-grade features for scalability and reproducibility. It's ready to use on HPC or AWS as you requested!